

Exercise 1 - List, Tuples, array, and NumPy

Answer the following questions:

1. What is the advantage of using Lists vs. Tuples

Lists are dynamic arrays that let us modify and resize the data we are storing, while **Tuples** are static arrays whose contents are fixed and immutable. This means that once a tuple is created, unlike a list, it cannot be modified or resized.

2. What is the advantage of using the array module vs. Python lists?

array modules have a static type and can store only that type of data, While type of elements in **Python lists** can be different, because lists only store 8-byte pointer to the actual object. But **array** modules is just a thin wrapper on C arrays. when storing same amount of data, **array** modules will use less space than **Python lists**.

array modules store data sequentially in memory, so that a slice of the array actually represents a continuous range in memory.

3. What are the memory fragmentation problem and Von-Neumann bottleneck? How do they affect the performance of a code? How can we try to address it?

Memory fragmentation problem: when our data is fragmented, we must move each piece over individually instead of moving the entire block over. This means we are invoking more memory transfer overhead, and we are forcing the CPU to wait while data is being transferred.

We can alleviate **Memory fragmentation problem** by **using the array module instead of lists**. And for any loop that does arithmetic on our array one element at a time to work on chunks of data.

Von-Neumann bottleneck: This refers to the limited bandwidth that exists between the memory and the CPU as a result of the tiered memory architecture that modern computers use.

To address **Von-Neumann bottleneck**, CPU try to predict the next instruction and load the relevant portions of memory into the cache while still working on the current instruction. And the best way to minimize the effects of the bottleneck is to be smart about how we allocate our memory and how we do our calculations over our data.

4. What is a page fault? What is the difference between a minor and a major page fault?

- 1) A **page-fault** is part of the modern memory allocation scheme.
- 2) When memory is first used, the OS throws a **minor page fault** interrupt, which pauses the program that is being run and properly allocates the memory; **Major page fault** happens when the program requests data from a device(disk, network, etc.) that hasn't been read yet. These are

even more expensive operations: not only do they interrupt your program, but they also involve reading from whichever device the data lives on.

5. What is the impact of a cache miss on the performance?

Cache misses can be a source of slowdowns, since we need to wait to fetch the data from RAM and we interrupt the flow of our execution pipeline

6. Which HPC libraries does your NumPy installation use? *Hint: you can check by writing a simple code.*

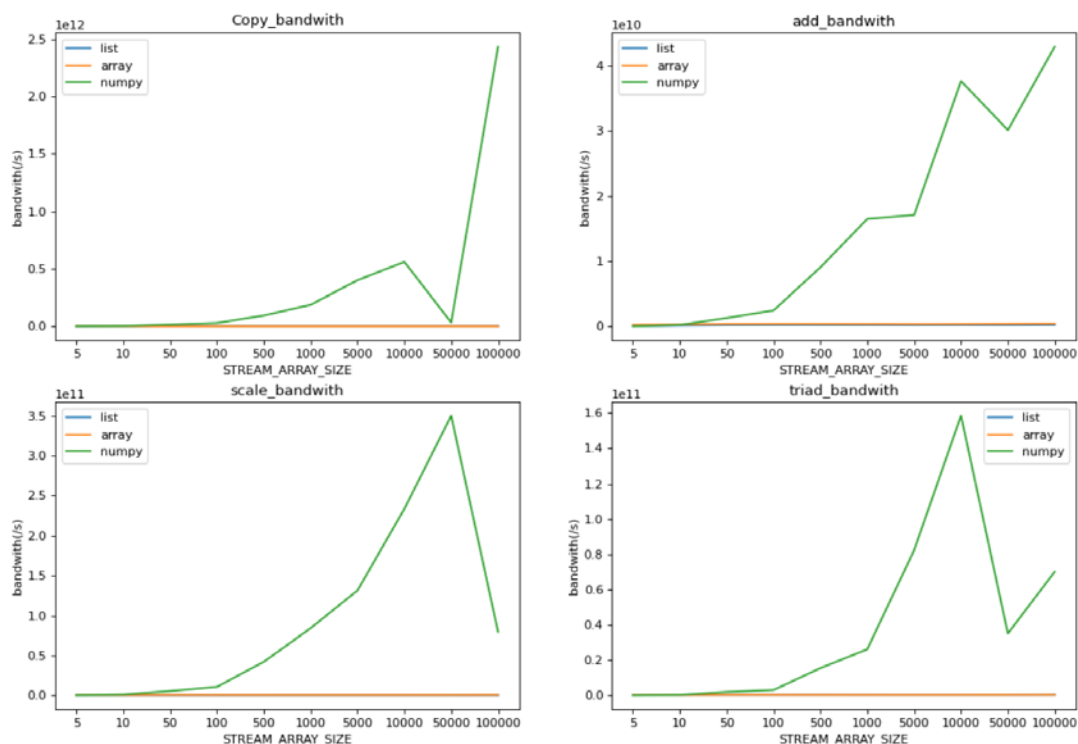
```
(base) PS C:\Users\daybeha> python
Python 3.8.8 (default, Apr 13 2021, 15:08:03) [MSC v.1916 64 bit (AMD64)] :: Anaconda, Inc. on win32
Type 'help', 'copyright', 'credits' or 'license' for more information.
>>> import numpy as np
>>> np.show_config()
openblas64_info:
  library_dirs = ['D:\\a\\numpy\\numpy\\build\\openblas64_info']
  libraries = ['openblas64_info']
  language = f77
  define_macros = [('HAVE_CBLAS', None), ('BLAS_SYMBOL_SUFFIX', '64_'), ('HAVE_BLAS_ILP64', None)]
blas_ilp64_opt_info:
  library_dirs = ['D:\\a\\numpy\\numpy\\build\\openblas64_info']
  libraries = ['openblas64_info']
  language = f77
  define_macros = [('HAVE_CBLAS', None), ('BLAS_SYMBOL_SUFFIX', '64_'), ('HAVE_BLAS_ILP64', None)]
openblas64_lapack_info:
  library_dirs = ['D:\\a\\numpy\\numpy\\build\\openblas64_lapack_info']
  libraries = ['openblas64_lapack_info']
  language = f77
  define_macros = [('HAVE_CBLAS', None), ('BLAS_SYMBOL_SUFFIX', '64_'), ('HAVE_BLAS_ILP64', None), ('HAVE_LAPACK', None)]
lapack_ilp64_opt_info:
  library_dirs = ['D:\\a\\numpy\\numpy\\build\\openblas64_lapack_info']
  libraries = ['openblas64_lapack_info']
  language = f77
  define_macros = [('HAVE_CBLAS', None), ('BLAS_SYMBOL_SUFFIX', '64_'), ('HAVE_BLAS_ILP64', None), ('HAVE_LAPACK', None)]
Supported SIMD extensions in this NumPy install:
  baseline = SSE, SSE2, SSE3
  found = SSSE3, SSE41, POPCNT, SSE42, AVX, F16C, FMA3, AVX2
  not found = AVX512F, AVX512CD, AVX512_SKX, AVX512_CLX, AVX512_CNL, AVX512_ICL
```

OpenBLAS64

Exercise 2 - STREAM Benchmark in Python to Measure the Memory Bandwidth

Task 2.1 Implement in Python the STREAM benchmark using Python lists, arrays from the array module, and NumPy arrays.

Task 2.2 Measure the bandwidth for the three Python array implementations (lists, array and numpy) varying the STREAM_ARRAY_SIZE and plot the results. Answer the questions: How does the bandwidth vary when increasing the STREAM_ARRAY_SIZE and why? How do the different implementation bandwidths compare to each other?



As shown above, bandwidth of list and array are very close, and bandwidth of numpy is lowest when STREAM_ARRAY_SIZE is low while far beyond list and array as STREAM_ARRAY_SIZE grows. The beginning is lowest is because numpy need to transfer values to object which increases the expense. But as STREAM_ARRAY_SIZE grows, the advantage of contiguous memory in numpy show up.

Exercise 3 - PyTest with the Julia Set Code

Task 3.1 Implement a separate code to test the assertion above using the pytest framework.

Code:

```
import pytest
@pytest.mark.parametrize('num1, num2, expected', [(1000, 300,
33219980),
(100, 300, 334236), (100, 100, 131532), (180, 300, 1076586), (180,
200, 1076586)])
def test_calc_pure_python(num1, num2, expected):
    assert sum(calc_pure_python(desired_width=num1,
max_iterations=num2)) == expected
```

Run: `pytest .\JuliaSet.py`

Output:

```
===== test session starts =====
collecting ... collected 5 items

JuliaSet.py::test_calc_pure_python[1000-300-33219980]
JuliaSet.py::test_calc_pure_python[100-300-334236]
JuliaSet.py::test_calc_pure_python[100-100-131532]
JuliaSet.py::test_calc_pure_python[180-300-1076586]
JuliaSet.py::test_calc_pure_python[180-200-1076586] PASSED [ 20%]calculate_z_serial_purepython took 5.472217798233832 seconds
PASSED [ 40%]calculate_z_serial_purepython took 0.05482816696166992 seconds
PASSED [ 60%]calculate_z_serial_purepython took 0.023900240325927734 seconds
PASSED [ 80%]calculate_z_serial_purepython took 0.18157362937927246 seconds
FAILED [100%]calculate_z_serial_purepython took 0.12715681921081543 seconds

JuliaSet.py:70 (test_calc_pure_python[180-200-1076586])
751254 != 1076586

Expected :1076586
Actual   :751254
<Click to see difference>

num1 = 180, num2 = 200, expected = 1076586

@pytest.mark.parametrize('num1, num2, expected', [(1000, 300, 33219980),
(100, 300, 334236), (100, 100, 131532), (180, 300, 1076586), (180, 200, 1076586)])
def test_calc_pure_python(num1, num2, expected):
    > assert sum(calc_pure_python(desired_width=num1, max_iterations=num2)) == expected
E     assert 751254 == 1076586
E       + where 751254 = sum([0, 0, 0, 0, 0, 0, ...])
E       +   where [0, 0, 0, 0, 0, 0, ...] = calc_pure_python(desired_width=180, max_iterations=200)

JuliaSet.py:74: AssertionError

===== 1 failed, 4 passed in 6.28s =====
```

Task 3.2 How would you implement the unit test with the possibility of having a different number of iterations and grid points? Implementation is optional.

If we know the result, we can just test by it. If we don't, we can calculate the shape of output as usual. So we can test by the output shape.

Exercise 4 - Python DGEMM Benchmark Operation

Answer the following questions:

- For which kind of problems do you use the BLAS libraries ?

When we need to calculate a large amount of number or matrix, and have requirements of time consuming.

- What is the difference between BLAS level-1, level-2 and level-3?

BLAS level-1 were limited to vector operations;

Level-2 provide routines for matrix-vector;

Level-3 provide routines for matrix-matrix;

Task 4.1 Implement the DGEMM with matrices as NumPy array

Code:

```
import numpy as np

N = 5

A = np.random.random((N, N)).astype(np.float64)
B = np.random.random((N, N)).astype(np.float64)
print(f"A:\n{A}\nB:\n{B}\n")

C = np.zeros_like(A)
# Multiplying first and second matrices and storing it in result
for i in range(N):
    for j in range(N):
        for k in range(N):
            C[i][j] += A[i][k] * B[k][j]

print(f"Result:\n{C}")
```

Output:

```

A:
[[0.45427996 0.93410411 0.43667526 0.51488453 0.4109359 ]
 [0.7041321  0.78952869 0.06591291 0.93581901 0.70911714]
 [0.60446704 0.94156487 0.92020224 0.0183316  0.60946608]
 [0.17245395 0.7982487  0.33645573 0.65606593 0.68943622]
 [0.07190902 0.57348201 0.18412812 0.77623531 0.06461922]]
B[[0.43728348 0.09022857 0.37619573 0.01525552 0.38051541]
 [0.99077245 0.0313196  0.7873943  0.957936  0.52984774]
 [0.72182833 0.19059772 0.62299358 0.08354924 0.93917315]
 [0.27000756 0.8186325  0.50687042 0.84327751 0.00387851]
 [0.72117571 0.48371393 0.11740592 0.10319876 0.73022721]]

Result:
[[1.87471803 0.77375075 1.48767837 1.41482474 1.37998072]
 [1.90180268 1.2099251  1.48521873 1.63490197 1.26961295]
 [2.30590983 0.56923216 1.62290749 1.06641731 2.03824409]
 [1.78350479 0.9752556  1.31650699 1.41980727 1.31055149]
 [0.98873492 0.72625254 0.99435637 1.22709024 0.55434617]]

```

Task 4.2 Using pytest develop a unit test for checking the correctness of your implementations.

Code:

```

import numpy as np

N = 5
A = np.ones((N,N)).astype(np.float64)
B = np.ones((N,N)).astype(np.float64)
print(f"A:\n{A}\nB:\n{B}\n")

def dgemm(A, B):
    C = np.zeros_like(A)
    # Multiplying first and second matrices and storing it in result
    for i in range(N):
        for j in range(N):
            for k in range(N):
                C[i][j] += A[i][k] * B[k][j]

    print(f"Result:\n{C}")
    return C

C = dgemm(A,B)

import pytest
def test_dgemm():
    assert (C==np.ones((N,N))*5).all()

```

Output:

```
===== test session starts =====
collecting ... collected 1 item

e_4_dgemm.py::test_dgemm PASSED [100%]

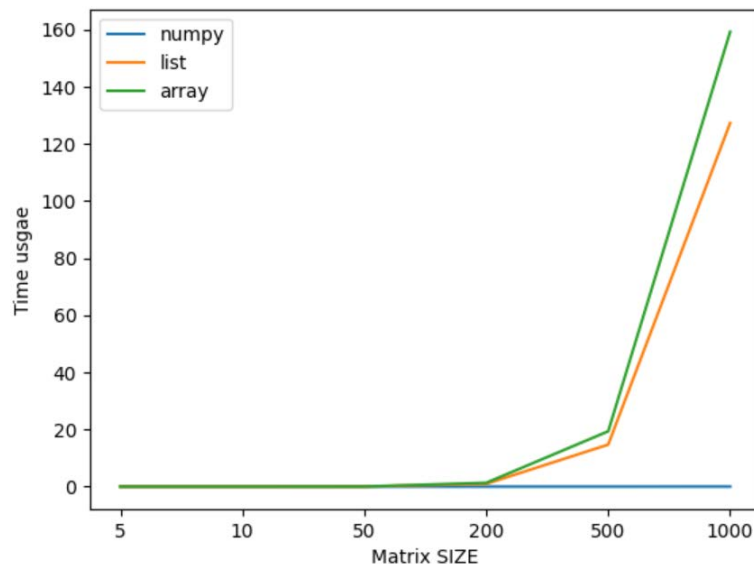
===== 1 passed in 0.02s =====
```

Task 4.2 Measure the execution time for each approach varying the size of the matrix. Report the average and error (std. deviation or min/max or interval of confidence). Answer the question: how the computational performance varies with increasing the size of the matrices and why so?

Task 4.3 Using the timing information and the number of operations for the DGEMM, calculate the FLOPS/s. How many operations are carried out in DGEMM with N as the matrices dimension? *Hint:* Think about the number of iterations completed in the loops and the number of flops per iteration. How do the FLOPS/s you measured compares to the theoretical peak of your processor (if we assume that we do one operation per cycle, then the peak is the clock frequency value)

Set N as [5,10,50,200,500,1000]

```
Numpy time   average:22.636333333328853 ms   min:92.000000000042508 us   max:0.010604299999997124 s
list time    average:235937.79383333333 ms   min:220.99999999913854 us   max:126.4852885 s
array time   average:302671.2825 ms   min:261.9999999997624 us   max:161.37949329999998 s
```



Answer4.2: As the size of matrices increase, the time usage non-linear increases in list and array. That is because the memory can no longer hold the values contiguous, which increases cache-miss.

Answer4.3: The number of operations in DGEMM with N as the matrices dimension is

$$N^2 * 2N = 2N^3$$

-----FLOPS/s-----				
		numpy	list	array
5		347801.8920422833	11210762.331823166	9363295.880133634
10		114942528.73526055	14760147.601480905	12070006.035007726
50		6925207756.235878	16134236.850596929	11936535.826318618
200		19316672703.126793	16668699.900789984	12035215.944494786
500		110894251242.01611	17001932.650489017	12903995.250875525
1000		192294749391.77063	15701986.705879984	12553818.25022325

The clock frequency of my CPU is 3.2GHz= $3.2 * 2^{30} = 3,435,973,836.8$

It can be seen that as N grows, the FLOPS/s of numpy grows quickly and far beyond the theoretical peak of my processor if we assume one operation per cycle.

Exercise 5 - A Python Discrete Fourier Transform

Task 5.1 Develop a DFT in Python and a unit test with pytest to check the calculation's correctness. Also, use the Python logging module to log the results. The data structures (lists, array, or NumPy) are of your choice.

Code:

```
N = 1024 # 采样点数
sample_freq = 120 # 采样频率 120 Hz, 大于两倍的最高频率
# sample_interval = 1 / sample_freq # 采样间隔
signal_len = N / sample_freq # 信号长度
t = np.arange(0, signal_len, 1 / sample_freq)

signal = 3 * np.sin(2 * np.pi * 20 * t) # 采集的信号
def DFT(xr, xi):
    N = len(X)
    X = np.zeros(N, np.complex_)
    Xr_o = X.real
    Xi_o = X.imag
    for k in range(N):
        for n in range(N):
            # Real part of X[k]
            Xr_o[k] += xr[n] * np.cos(n * k * 2 * np.pi / N) + xi[n] *
np.sin(n * k * 2 * np.pi / N)
            # Imaginary part of X[k]
            Xi_o[k] += -xr[n] * np.sin(n * k * 2 * np.pi / N) + xi[n] *
np.cos(n * k * 2 * np.pi / N)

    return X

import pytest
def test_DFT():
    Freq = np.zeros(N, np.complex_)
    Freq = DFT(signal, np.zeros_like(signal))

    fft_data = fft(signal)
    assert ((Freq-fft_data)<10e-6).all()
```

```

===== test session starts =====
collecting ... collected 1 item

e5_DFT.py::test_DFT PASSED [100%]

===== 1 passed in 4.71s =====

```

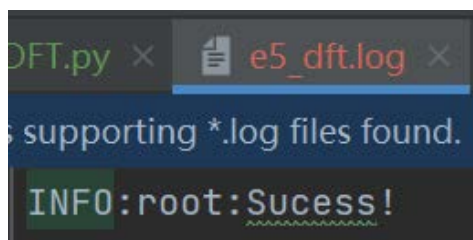
```

import logging
logging.basicConfig(filename='e5_dft.log', level=logging.INFO)

Freq = np.zeros(N, np.complex_)
Freq = DFT(signal, np.zeros_like(signal), Freq)
fft_data = fft(signal)
logging.info("Sucess!")

if ((Freq-fft_data)<10e-6).all():
    logging.info("Sucess!")
else:
    logging.error("Something went wrong")

```



Task 5.2 Document the code using docstrings and generate automatic HTML documentation.

Assginment2

Navigation

Quick search

 Go

Hide Search Matches

Welcome to Assginment2's documentation!

Indices and tables

- [Index](#)
- [Module Index](#)
- [Search Page](#)

©2023, TengfeiLu. | Powered by Sphinx 6.1.3 & Alabaster 0.7.13 | [Page source](#)

Task 5.3 Check your code using a Python linter. Address the issues raised by the linter, including the style issues. To address some of the issues, you can use a Python auto-formatter.

```

(evaluator) PS D:\win\桌面\KTH\courses\DD2358 Introduction to High Performance Computing\Assignments\Assignment 2> pylint .\e5_DFT.py
***** Module e5_DFT
e5_DFT.py:1:0: C0103: Module name "e5_DFT" doesn't conform to snake_case naming style (invalid-name)
e5_DFT.py:44:0: C0103: Function name "DFT" doesn't conform to snake_case naming style (invalid-name)
e5_DFT.py:60:4: C0103: Variable name "X" doesn't conform to snake_case naming style (invalid-name)
e5_DFT.py:61:4: C0103: Variable name "X_r" doesn't conform to snake_case naming style (invalid-name)
e5_DFT.py:62:4: C0103: Variable name "X_i" doesn't conform to snake_case naming style (invalid-name)
e5_DFT.py:64:12: C0103: Variable name "n" doesn't conform to snake_case naming style (invalid-name)

Your code has been rated at 8.12/10 (previous run: 8.12/10, +0.00)

```


Task 5.4 Measure the execution time, varying the input size from 8 to 1024 elements, and plot it.

Code:

```
times = []
Ns = [8,16,32,64,128,256,512,1024]
for N in Ns:
    print(f"N:{N}")
    t = timer()

    Freq = np.zeros(N, np.complex_)
    Freq = DFT(signal, np.zeros_like(signal), Freq)
    times.append(timer()-t)
plt.plot(range(len(Ns)), times)
plt.xticks(range(len(Ns)), labels=Ns)
plt.xlabel("Input SIZE(N)")
plt.ylabel("Time Usage(s)")
plt.show()
```

Output:



Task 5.5 Profile the code with all the profiling tools that can be useful for performance analysis (from coarse-grained to fine-grained), fixing the input size, e.g., 1024. Motivate the choice of profiling tools and report the profiling results

Considering there are only DFT() functions used, just line_profiler instead of cProfile.

line_profiler:

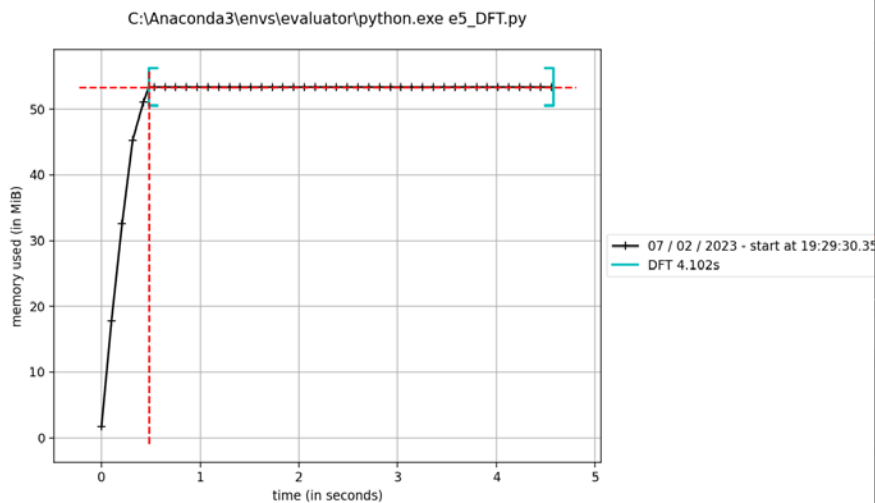
Total time: 5.63189 s
File: .\e5_DFT.py
Function: DFT at line 41

Line #	Hits	Time	Per Hit	% Time	Line Contents
41					@profile
42					def DFT(x_r, x_i):
43					"""
44					DFT calculator
45					
46					an implementation of DFT calculator
47					
48					Parameters:
49					
50					xr (np.array): real part of signal
51					
52					xi (np.array): image part of signal
53					
54					Returns:
55					
56					X(np.array): DFT result
57					"""
58	1	8.1	8.1	0.0	X = np.zeros(N, np.complex_)
59	1	1.7	1.7	0.0	X_r = X.real
60	1	1.4	1.4	0.0	X_i = X.imag
61	1024	222.4	0.2	0.0	for k in range(N):
62	1048576	220338.1	0.2	3.9	for n in range(N):
63					# Real part of X[k]
64	1048576	1389639.9	1.3	24.7	X_r[k] += x_r[n] * np.cos(n * k * 2 * np.pi / N) \
65	1048576	1314279.8	1.3	23.3	+ x_i[n] * np.sin(n * k * 2 * np.pi / N)
66					# Imaginary part of X[k]
67	1048576	1411542.7	1.3	25.1	X_i[k] += -x_r[n] * np.sin(n * k * 2 * np.pi / N) \
68	1048576	1295853.3	1.2	23.0	+ x_i[n] * np.cos(n * k * 2 * np.pi / N)
69	1	0.6	0.6	0.0	return X

Memory_profiler

Line #	Mem usage	Increment	Occurrences	Line Contents
41	53.844 MiB	53.844 MiB	1	@profile
42				def DFT(x_r, x_i):
43				"""
44				DFT calculator
45				
46				an implementation of DFT calculator
47				
48				Parameters:
49				
50				xr (np.array): real part of signal
51				
52				xi (np.array): image part of signal
53				
54				Returns:
55				
56				X(np.array): DFT result
57				"""
58	53.848 MiB	0.004 MiB	1	X = np.zeros(N, np.complex_)
59	53.848 MiB	0.000 MiB	1	X_r = X.real
60	53.848 MiB	0.000 MiB	1	X_i = X.imag
61	53.898 MiB	-26.086 MiB	1025	for k in range(N):
62	53.898 MiB	-26698.652 MiB	1049600	for n in range(N):
63				# Real part of X[k]
64	53.898 MiB	-53345.207 MiB	2097152	X_r[k] += x_r[n] * np.cos(n * k * 2 * np.pi / N) \
65	53.898 MiB	-26672.625 MiB	1048576	+ x_i[n] * np.sin(n * k * 2 * np.pi / N)
66				# Imaginary part of X[k]
67	53.898 MiB	-53345.223 MiB	2097152	X_i[k] += -x_r[n] * np.sin(n * k * 2 * np.pi / N) \
68	53.898 MiB	-26672.625 MiB	1048576	+ x_i[n] * np.cos(n * k * 2 * np.pi / N)
69	53.844 MiB	-0.055 MiB	1	return X

mprof



Exercise 6 - Experiment with the Python Debugger

As part of this exercise, we ask you to complete an online tutorial on the Python pdb debugger. Follow the instructions at <https://github.com/spide/pdb-tutorial>. [Links to an external site.](#)

Task 6.1 Reflection: answer the questions: What are the advantages of using a debugger? What challenges did you find in using the pdb debugger, if any?

1) With a debugger, we can:

- Explore the state of a running program
- Test implementation code before applying it
- Follow the program's execution logic

2) pdb debugger will make our code bloat. Every IDEA of Python , like Pycharm, VSCode etc. , have the interface for debug and its easier to use. Except that, IDEAs have many other features like autocomplete which greatly improve programming process.

Bonus Exercise - Performance Analysis and Optimization of the Game of Life Code

Task B.1 Check the code with a linter, and in case, run an auto-formatter. Produce HTML documentation running sphinx.

Before auto-formatter:

```
pylint .\conway.py
***** Module conway
conway.py:84:0: R1707: Disallow trailing comma tuple (trailing-comma-tuple)
conway.py:142:0: C0304: Final newline missing (missing-final-newline)
conway.py:9:0: C0410: Multiple imports on one line (sys, argparse) (multiple-imports)
conway.py:12:0: R0402: Use 'from matplotlib import animation' instead (consider-using-from-import)
conway.py:19:0: C0103: Function name 'randomGrid' doesn't conform to snake_case naming style (invalid-name)
conway.py:19:15: C0103: Argument name 'N' doesn't conform to snake_case naming style (invalid-name)
conway.py:24:0: C0103: Function name 'addGlider' doesn't conform to snake_case naming style (invalid-name)
conway.py:32:0: C0103: Function name 'addGosperGliderGun' doesn't conform to snake_case naming style (invalid-name)

conway.py:61:0: C0116: Missing function or method docstring (missing-function-docstring)
conway.py:61:11: C0103: Argument name 'frameNum' doesn't conform to snake_case naming style (invalid-name)
conway.py:61:32: C0103: Argument name 'N' doesn't conform to snake_case naming style (invalid-name)
conway.py:64:4: C0103: Variable name 'newGrid' doesn't conform to snake_case naming style (invalid-name)
conway.py:61:11: W0613: Unused argument 'frameNum' (unused-argument)
conway.py:88:0: C0116: Missing function or method docstring (missing-function-docstring)
conway.py:102:4: C0103: Variable name 'N' doesn't conform to snake_case naming style (invalid-name)
conway.py:104:8: C0103: Variable name 'N' doesn't conform to snake_case naming style (invalid-name)
conway.py:107:4: C0103: Variable name 'updateInterval' doesn't conform to snake_case naming style (invalid-name)
conway.py:109:8: C0103: Variable name 'updateInterval' doesn't conform to snake_case naming style (invalid-name)
conway.py:125:9: C0103: Variable name 'ax' doesn't conform to snake_case naming style (invalid-name)
conway.py:9:0: W0611: Unused import sys (unused-import)

Your code has been rated at 7.37/10
```

After auto-formatter:

```
black .\conway.py
reformatted conway.py

All done! ✨ 🍰 ✨
1 file reformatted.
(evaluator) PS D:\win\桌面\KTH\courses\DD2358 Introduction to High Performance Computing\Assignments\Assignment 2>
pylint .\conway.py
***** Module conway
conway.py:9:0: C0410: Multiple imports on one line (sys, argparse) (multiple-imports)
conway.py:12:0: R0402: Use 'from matplotlib import animation' instead (consider-using-from-import)
conway.py:19:0: C0103: Function name 'randomGrid' doesn't conform to snake_case naming style (invalid-name)
conway.py:19:15: C0103: Argument name 'N' doesn't conform to snake_case naming style (invalid-name)
conway.py:24:0: C0103: Function name 'addGlider' doesn't conform to snake_case naming style (invalid-name)
conway.py:30:0: C0103: Function name 'addGosperGliderGun' doesn't conform to snake_case naming style (invalid-name)

conway.py:59:0: C0116: Missing function or method docstring (missing-function-docstring)
conway.py:59:11: C0103: Argument name 'frameNum' doesn't conform to snake_case naming style (invalid-name)
conway.py:59:32: C0103: Argument name 'N' doesn't conform to snake_case naming style (invalid-name)
conway.py:62:4: C0103: Variable name 'newGrid' doesn't conform to snake_case naming style (invalid-name)
conway.py:59:11: W0613: Unused argument 'frameNum' (unused-argument)
conway.py:95:0: C0116: Missing function or method docstring (missing-function-docstring)
conway.py:111:4: C0103: Variable name 'N' doesn't conform to snake_case naming style (invalid-name)
conway.py:113:8: C0103: Variable name 'N' doesn't conform to snake_case naming style (invalid-name)
conway.py:116:4: C0103: Variable name 'updateInterval' doesn't conform to snake_case naming style (invalid-name)
conway.py:118:8: C0103: Variable name 'updateInterval' doesn't conform to snake_case naming style (invalid-name)
conway.py:134:9: C0103: Variable name 'ax' doesn't conform to snake_case naming style (invalid-name)
conway.py:9:0: W0611: Unused import sys (unused-import)

Your code has been rated at 7.63/10 (previous run: 7.37/10, +0.26)
```

Sphinx document

conway

Navigation

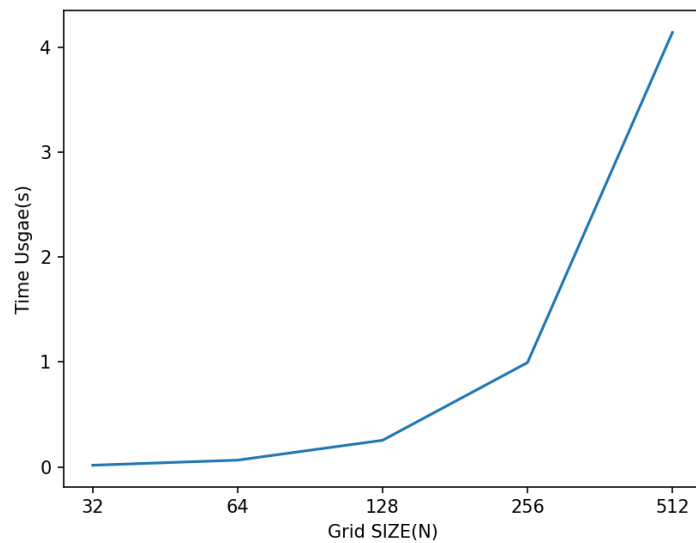
Quick search

Go

Welcome to conway's
documentation!
Indices and tables

- [Index](#)
- [Module Index](#)
- [Search Page](#)

Task B.2 Measure the execution time, varying the grid size (and fixed number of iterations). Make a plot with this information.



Task B.3 Use different profilers (from coarse- to fine-grained) to identify performance bottlenecks and potential improvement. Report the results of the profilers. The choice of profilers is up to you.

line_profiler:

```
Total time: 17.7194 s
File: .\conway.py
Function: update_ at line 95
```

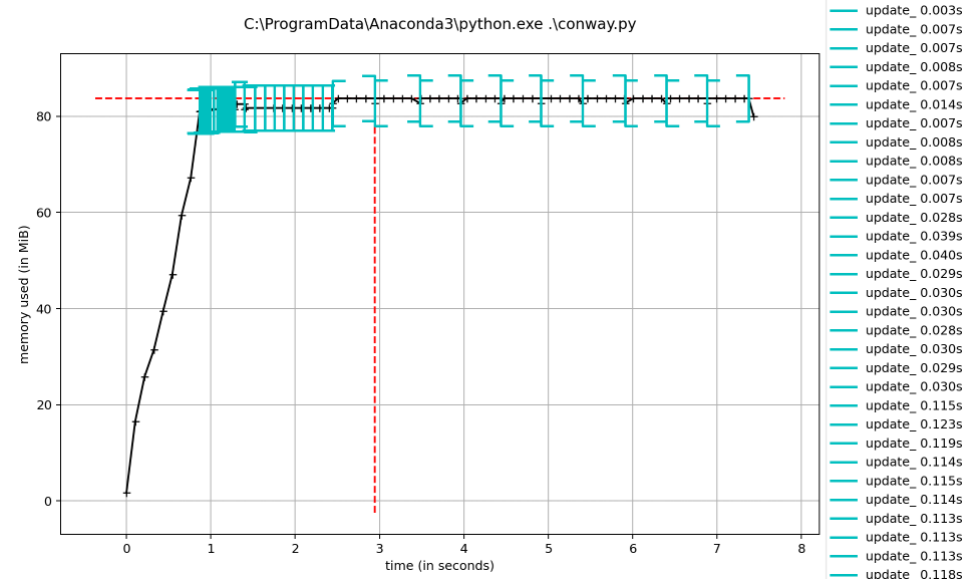
Line #	Hits	Time	Per Hit	% Time	Line Contents
95					@profile
96					def update_(grid, N):
97					# copy grid since we require 8 neighbors for calculation
98					# and we go line by line
99	50	2810.4	56.2	0.0	newGrid = grid.copy()
100	9920	2019.1	0.2	0.0	for i in range(N):
101	3491840	908259.9	0.3	5.1	for j in range(N):
102					# compute 8-neghbor sum
103					# using toroidal boundary conditions - x and y wrap around
104					# so that the simulation takes place on a toroidal surface.
105	3491840	1065499.5	0.3	6.0	total = int(
106	3491840	1119548.5	0.3	6.3	(
107	3491840	1754562.0	0.5	9.9	grid[i, (j - 1) % N]
108	3491840	1334396.6	0.4	7.5	+ grid[i, (j + 1) % N]
109	3491840	1358046.3	0.4	7.7	+ grid[(i - 1) % N, j]
110	3491840	1285906.6	0.4	7.3	+ grid[(i + 1) % N, j]
111	3491840	1529260.2	0.4	8.6	+ grid[(i - 1) % N, (j - 1) % N]
112	3491840	1473136.2	0.4	8.3	+ grid[(i - 1) % N, (j + 1) % N]
113	3491840	1523843.2	0.4	8.6	+ grid[(i + 1) % N, (j - 1) % N]
114	3491840	1448156.6	0.4	8.2	+ grid[(i + 1) % N, (j + 1) % N]
115)
116	3491840	647310.4	0.2	3.7	/ 255
117)
118					# apply Conway's rules
119	2870196	1412281.5	0.5	8.0	if grid[i, j] == ON:
120	331862	101111.4	0.3	0.6	if (total < 2) or (total > 3):
121	289782	101067.5	0.3	0.6	newGrid[i, j] = OFF
122					else:
123	2594374	555058.6	0.2	3.1	if total == 3:
124	275822	97144.3	0.4	0.5	newGrid[i, j] = ON
125	50	18.2	0.4	0.0	return newGrid

Memory_profiler

Filename: .\conway.py

Line #	Mem usage	Increment	Occurrences	Line Contents
95	82.742 MiB	66.234 MiB	50	@profile
96				def update_grid(N):
97				# copy grid since we require 8 neighbors for calculation
98				# and we go line by line
99	83.746 MiB	3.023 MiB	50	newGrid = grid.copy()
100	83.758 MiB	-1871.859 MiB	9970	for i in range(N):
101	83.758 MiB	-502326.684 MiB	3501760	for j in range(N):
102				# compute 8-neighbor sum
103				# using toroidal boundary conditions - x and y wrap around
104				# so that the simulation takes place on a toroidal surface.
105	83.758 MiB	-1000923.875 MiB	6983680	total = int(
106	83.758 MiB	-500461.973 MiB	3491840	(
107	83.758 MiB	-4003695.320 MiB	27934720	grid[i, (j - 1) % N]
108	83.758 MiB	-500461.875 MiB	3491840	+ grid[i, (j + 1) % N]
109	83.758 MiB	-500461.875 MiB	3491840	+ grid[(i - 1) % N, j]
110	83.758 MiB	-500461.902 MiB	3491840	+ grid[(i + 1) % N, j]
111	83.758 MiB	-500461.922 MiB	3491840	+ grid[(i - 1) % N, (j - 1) % N]
112	83.758 MiB	-500461.938 MiB	3491840	+ grid[(i - 1) % N, (j + 1) % N]
113	83.758 MiB	-500461.949 MiB	3491840	+ grid[(i + 1) % N, (j - 1) % N]
114	83.758 MiB	-500461.941 MiB	3491840	+ grid[(i + 1) % N, (j + 1) % N]
115)
116	83.758 MiB	-500461.938 MiB	3491840	/ 255
117)
118				# apply Conway's rules
119	83.758 MiB	-500462.000 MiB	3491840	if grid[i, j] == ON:
120	83.758 MiB	-86716.445 MiB	612111	if (total < 2) or (total > 3):
121	83.758 MiB	-39303.988 MiB	285268	newGrid[i, j] = OFF
122				else:
123	83.758 MiB	-413745.555 MiB	2879729	if total == 3:
124	83.758 MiB	-36779.961 MiB	270364	newGrid[i, j] = ON
125	83.746 MiB	-7.234 MiB	50	return newGrid

mprof



Task B.4 Implement an optimization, report the new profiling results and show the performance improvement.

Updated Code:

```
def update_grid(N):
    # copy grid since we require 8 neighbors for calculation
    # and we go line by line
    newGrid = grid.copy()
    totals = np.zeros_like(grid)
```



```


    for i in range(N):
        for j in range(N):
            if i in [0, N - 1] and j in [0, N - 1]:
                totals[i, j] = int((grid[i, (j - 1) % N] + grid[i, (j +
1) % N] +
                                grid[(i - 1) % N, j] + grid[(i + 1) %
N, j] +
                                grid[(i - 1) % N, (j - 1) % N] +
grid[(i - 1) % N, (j + 1) % N] +
                                grid[(i + 1) % N, (j - 1) % N] +
grid[(i + 1) % N, (j + 1) % N]) / 255)
                totals[1:-1, 1:-1] += ((grid[:-2, :-2] + grid[:-2, 1:-1] +
grid[:-2, 2:] +
                                grid[1:-1, :-2] + grid[1:-1, 2:] +
                                grid[2:, :-2] + grid[2:, 1:-1] + grid[2:,
2:])/255).astype(np.int32)

            newGrid[np.logical_and(grid == ON, np.logical_or((totals < 2),
(totals > 3)))] = OFF
            newGrid[np.logical_and(grid != ON, totals == 3)] = ON

    return newGrid

```

After update:



line_profiler (greatly optimized)

```

Timer unit: 1e-06 s

Total time: 1.62476 s
File: .\conway.py
Function: update_ at line 98

Line #      Hits          Time Per Hit   % Time  Line Contents
=====
98          @profile
99      def update_(grid, N):
100          # copy grid since we require 8 neighbors for calculation
101          # and we go line by line
102          newGrid = grid.copy()
103          totals = np.zeros_like(grid)
104
105          for i in range(N):
106              for j in range(N):
107                  if i in [0, N - 1] and j in [0, N - 1]:
108                      totals[i, j] = int((grid[i, (j - 1) % N] + grid[i, (j + 1) % N] +
109                      grid[(i - 1) % N, j] + grid[(i + 1) % N, j] +
110                      grid[(i - 1) % N, (j - 1) % N] + grid[(i - 1) % N, (j + 1) % N] +
111                      grid[(i + 1) % N, (j - 1) % N] + grid[(i + 1) % N, (j + 1) % N]) / 255)
112          totals[1:-1, 1:-1] += ((grid[:-2, :-2] + grid[:-2, 1:-1] + grid[:-2, 2:] +
113          grid[1:-1, :-2] + grid[1:-1, 2:] +
114          grid[2:, :-2] + grid[2:, 1:-1] + grid[2:, 2:])/255).astype(np.int32)
115
116          newGrid[np.logical_and(grid == ON, np.logical_or((totals < 2), (totals > 3)))] = OFF
117          newGrid[np.logical_and(grid != ON, totals == 3)] = ON
118
119          return newGrid

```

memory_profiler(on optimization)

```

Filename: .\conway.py

Line #      Mem usage      Increment   Occurrences   Line Contents
=====
98      83.676 MiB      55.859 MiB          50      @profile
99
100          # copy grid since we require 8 neighbors for calculation
101          # and we go line by line
102      84.680 MiB       2.656 MiB          50      newGrid = grid.copy()
103      85.684 MiB       4.711 MiB          50      totals = np.zeros_like(grid)
104
105      85.684 MiB     -1463.879 MiB        9970      for i in range(N):
106      85.684 MiB     -416149.648 MiB      3501760      for j in range(N):
107      85.684 MiB     -414691.164 MiB      3491840      if i in [0, N - 1] and j in [0, N - 1]:
108      85.684 MiB     -173.117 MiB        1600      totals[i, j] = int((grid[i, (j - 1) % N] + grid[i, (j + 1) % N] +
109      85.684 MiB      -43.281 MiB          400      grid[(i - 1) % N, j] + grid[(i + 1) % N, j] +
110      85.684 MiB      -43.281 MiB          400      grid[(i - 1) % N, (j - 1) % N] + grid[(i - 1) % N, (j + 1) % N] +
111      85.684 MiB      -64.922 MiB          600      grid[(i + 1) % N, (j - 1) % N] + grid[(i + 1) % N, (j + 1) % N]) / 255)
112      86.922 MiB     -18.703 MiB          400      totals[1:-1, 1:-1] += ((grid[:-2, :-2] + grid[:-2, 1:-1] + grid[:-2, 2:] +
113      86.922 MiB     -15.250 MiB          100      grid[1:-1, :-2] + grid[1:-1, 2:] +
114      86.918 MiB     -11.105 MiB          250      grid[2:, :-2] + grid[2:, 1:-1] + grid[2:, 2:])/255).astype(np.int32)
115
116      84.930 MiB     -20.270 MiB          50      newGrid[np.logical_and(grid == ON, np.logical_or((totals < 2), (totals > 3)))] = OFF
117      85.684 MiB       7.168 MiB          50      newGrid[np.logical_and(grid != ON, totals == 3)] = ON
118
119      85.684 MiB      -0.379 MiB          50      return newGrid

```