

**GitHub:** [https://github.com/DayBeha/HPC\\_assignment](https://github.com/DayBeha/HPC_assignment)

## Exercise 1 - Cythonize the STREAM Benchmark

**Task 1.1:** Develop a Cython version of the STREAM benchmark. Make sure to define the *ctypes* for obtaining full performance.

**Cython Code as follows:**

```
#cython: language_level=3

import numpy as np
cimport numpy as np

def copy(float[:] a, float[:] b):
    # cdef float[:] c = np.empty(len(a), dtype=np.float32)
    for i in range(len(a)):
        b[i] = a[i]
    # return c

def scale(float scalar, float[:] a, float[:] b):
    # cdef float[:] b = np.empty(len(a), dtype=np.float32)
    for i in range(len(a)):
        b[i] = scalar * a[i]
    # return b

def add(float[:] a, float[:] b, float[:] c):
    # cdef float[:] c = np.empty(len(a), dtype=np.float32)
    for i in range(len(a)):
        c[i] = a[i] + b[i]
    # return c

def triad(float[:] a, float scalar, float[:] b, float[:] c):
    # cdef float[:] c = np.empty(len(a), dtype=np.float32)
    for i in range(len(a)):
        c[i] = a[i] + scalar * b[i]
    # return c
```

**Task 1.2:** Plot the bandwidth results varying the arrays' size. Answer the question: how does the bandwidth measured with Cython code compare to bandwidth obtained in Assignment II.

**bandwidth before Cython:**

STREAM\_ARRAY\_SIZE = 5

Copy	time use:8.00 us	data amount:280	bandwith:35.000 MB/s
add	time use:87.00 us	data amount:280	bandwith:3.218 MB/s
scale	time use:152.00 us	data amount:420	bandwith:2.763 MB/s

triad time use:31.00 us data amount:420 bandwidth:13.548 MB/s

STREAM\_ARRAY\_SIZE = 10

Copy time use:2.00 us data amount:560 bandwidth:280.000 MB/s

add time use:10.00 us data amount:560 bandwidth:56.000 MB/s

scale time use:21.00 us data amount:840 bandwidth:40.000 MB/s

triad time use:18.00 us data amount:840 bandwidth:46.667 MB/s

STREAM\_ARRAY\_SIZE = 50

Copy time use:1.00 us data amount:2800 bandwidth:2800.000 MB/s

add time use:8.00 us data amount:2800 bandwidth:350.000 MB/s

scale time use:16.00 us data amount:4200 bandwidth:262.500 MB/s

triad time use:18.00 us data amount:4200 bandwidth:233.333 MB/s

STREAM\_ARRAY\_SIZE = 100

Copy time use:2.00 us data amount:5600 bandwidth:2800.000 MB/s

add time use:8.00 us data amount:5600 bandwidth:700.000 MB/s

scale time use:18.00 us data amount:8400 bandwidth:466.667 MB/s

triad time use:21.00 us data amount:8400 bandwidth:400.000 MB/s

STREAM\_ARRAY\_SIZE = 500

Copy time use:4.00 us data amount:28000 bandwidth:7000.000 MB/s

add time use:10.00 us data amount:28000 bandwidth:2800.000 MB/s

scale time use:22.00 us data amount:42000 bandwidth:1909.091 MB/s

triad time use:25.00 us data amount:42000 bandwidth:1680.000 MB/s

STREAM\_ARRAY\_SIZE = 1000

Copy time use:4.00 us data amount:56000 bandwidth:14000.000 MB/s

add time use:12.00 us data amount:56000 bandwidth:4666.667 MB/s

scale time use:21.00 us data amount:84000 bandwidth:4000.000 MB/s

triad time use:25.00 us data amount:84000 bandwidth:3360.000 MB/s

STREAM\_ARRAY\_SIZE = 5000

Copy time use:4.00 us data amount:280000 bandwidth:70000.000 MB/s

add time use:17.00 us data amount:280000 bandwidth:16470.588 MB/s

scale time use:26.00 us data amount:420000 bandwidth:16153.846 MB/s

triad time use:91.00 us data amount:420000 bandwidth:4615.385 MB/s

STREAM\_ARRAY\_SIZE = 10000

Copy time use:4.00 us data amount:560000 bandwidth:140000.000 MB/s

add time use:36.00 us data amount:560000 bandwidth:15555.556 MB/s

scale time use:52.00 us data amount:840000 bandwidth:16153.846 MB/s


triad time use:60.00 us data amount:840000 bandwidth:14000.000 MB/s

STREAM\_ARRAY\_SIZE = 50000

Copy	time use:18.00 us	data amount:2800000	bandwith:155555.556 MB/s
add	time use:347.00 us	data amount:2800000	bandwith:8069.164 MB/s
scale	time use:242.00 us	data amount:4200000	bandwith:17355.372 MB/s
triad	time use:302.00 us	data amount:4200000	bandwith:13907.285 MB/s

STREAM\_ARRAY\_SIZE = 100000

Copy	time use:34.00 us	data amount:5600000	bandwith:164705.882 MB/s
add	time use:898.00 us	data amount:5600000	bandwith:6236.080 MB/s
scale	time use:1220.00 us	data amount:8400000	bandwith:6885.246 MB/s
triad	time use:1616.00 us	data amount:8400000	bandwith:5198.020 MB/s



### bandwidth after Cython:

STREAM\_ARRAY\_SIZE = 5

Copy	time use:83.00 us	data amount:280	bandwith:3.373 MB/s
add	time use:29.00 us	data amount:280	bandwith:9.655 MB/s
scale	time use:10.00 us	data amount:420	bandwith:42.000 MB/s
triad	time use:11.00 us	data amount:420	bandwith:38.182 MB/s

STREAM\_ARRAY\_SIZE = 10

Copy	time use:13.00 us	data amount:560	bandwith:43.077 MB/s
add	time use:11.00 us	data amount:560	bandwith:50.909 MB/s
scale	time use:8.00 us	data amount:840	bandwith:105.000 MB/s
triad	time use:9.00 us	data amount:840	bandwith:93.333 MB/s

STREAM\_ARRAY\_SIZE = 50

Copy      time use:9.00 us    data amount:2800    bandwidth:311.111 MB/s  
add        time use:9.00 us    data amount:2800    bandwidth:311.111 MB/s  
scale time use:7.00 us    data amount:4200    bandwidth:600.000 MB/s  
triad time use:9.00 us    data amount:4200    bandwidth:466.667 MB/s

STREAM\_ARRAY\_SIZE = 100

Copy      time use:9.00 us    data amount:5600    bandwidth:622.222 MB/s  
add        time use:10.00 us   data amount:5600    bandwidth:560.000 MB/s  
scale time use:7.00 us    data amount:8400    bandwidth:1200.000 MB/s  
triad time use:10.00 us   data amount:8400    bandwidth:840.000 MB/s

STREAM\_ARRAY\_SIZE = 500

Copy      time use:12.00 us   data amount:28000    bandwidth:2333.333 MB/s  
add        time use:15.00 us   data amount:28000    bandwidth:1866.667 MB/s  
scale time use:10.00 us   data amount:42000    bandwidth:4200.000 MB/s  
triad time use:15.00 us   data amount:42000    bandwidth:2800.000 MB/s

STREAM\_ARRAY\_SIZE = 1000

Copy      time use:17.00 us   data amount:56000    bandwidth:3294.118 MB/s  
add        time use:22.00 us   data amount:56000    bandwidth:2545.455 MB/s  
scale time use:14.00 us   data amount:84000    bandwidth:6000.000 MB/s  
triad time use:22.00 us   data amount:84000    bandwidth:3818.182 MB/s

STREAM\_ARRAY\_SIZE = 5000

Copy      time use:45.00 us   data amount:280000    bandwidth:6222.222 MB/s  
add        time use:74.00 us   data amount:280000    bandwidth:3783.784 MB/s  
scale time use:43.00 us   data amount:420000    bandwidth:9767.442 MB/s  
triad time use:74.00 us   data amount:420000    bandwidth:5675.676 MB/s

STREAM\_ARRAY\_SIZE = 10000

Copy      time use:82.00 us   data amount:560000    bandwidth:6829.268 MB/s  
add        time use:137.00 us   data amount:560000    bandwidth:4087.591 MB/s  
scale time use:79.00 us   data amount:840000    bandwidth:10632.911 MB/s  
triad time use:141.00 us   data amount:840000    bandwidth:5957.447 MB/s


STREAM\_ARRAY\_SIZE = 50000

Copy      time use:385.00 us   data amount:2800000    bandwidth:7272.727 MB/s  
add        time use:672.00 us   data amount:2800000    bandwidth:4166.667 MB/s  
scale time use:372.00 us   data amount:4200000    bandwidth:11290.323 MB/s  
triad time use:681.00 us   data amount:4200000    bandwidth:6167.401 MB/s

STREAM\_ARRAY\_SIZE = 100000

Copy      time use:749.00 us   data amount:5600000    bandwidth:7476.636 MB/s  
add        time use:1340.00 us   data amount:5600000    bandwidth:4179.104 MB/s

scale time use:741.00 us      data amount:8400000      bandwidth:11336.032 MB/s  
 triad time use:1355.00 us      data amount:8400000      bandwidth:6199.262 MB/s




## Conclusion:

- 1、As the plots shows, 'copy' get the highest bandwidth and high better than others before Cythonizing and 'scale' get the highest bandwidth after Cythonizing. It may because it's copied as matrix before Cythonizing.
- 2、The performance after Cythonizing is no better than before which not as expectation. I both tried pass variable to functions in \*.pyx and define a new array. No better performance appear.
- 3、The bandwidth decrease when the STREAM\_ARRAY\_SIZE is too large(100000) before Cythonizing but not decrease after Cythonizing where bandwidth after Cythonizing is better than before. It may because 'np.array' need more memory than 'cdef'.

## Exercise 2 - Gauss-Seidel for Poisson Solver

**Task 2.1:** Develop the Gauss-Seidel solver with Python List, array, or NumPy. Plot the performance varying the grid size.

Set grid size as [5,10,50,200,500,1000]



Numpy time	average:32650.33 ms	min:317.35 ms	max:154148.00 ms
list time	average:515427.62 ms	min:28.35 ms	max:2418517.64 ms
array time	average:628008.64 ms	min:33.97 ms	max:2949629.43 ms

-----FLOPS/s-----			
+-----+-----+-----+-----+			
	numpy	list	array
+-----+-----+-----+-----+			
5	7521.692561346932	88149.21899791942	73570.52470498308
10	63017.13435883211	122703.90321116058	99343.34051916863
50	6225316.619603266	470853.1538968467	389149.39414109127
200	134905431.29266375	1799314.2700859562	1471214.8944104011
500	63264281.550953686	4313136.6837794045	3556140.889391718
1000	129745408.30513468	8269528.148778352	6780512.704508332
+-----+-----+-----+-----+			

The performance is Numpy>list>array.

**Task 2.2:** Profile the code to identify the part of the code to optimize. You can use the tool of your choice.

### line\_profiler

```

Timer unit: 1e-06 s

Total time: 2.7237 s
File: .\e2GaussSeidel.py
Function: gauss_seidel at line 13

Line #      Hits          Time Per Hit   % Time  Line Contents
=====
    13                               @profile
    14                               def gauss_seidel(f):
    15      1000           718.3     0.7     0.0      newf = f.copy()
    16     48000          8986.6     0.2     0.3      for i in range(1, len(newf) - 1):
    17    2304000        440935.2     0.2    16.2          for j in range(1, len(newf[0]) - 1):
    18    4608000        1226510.6     0.3    45.0              newf[i][j] = 0.25 * (newf[i][j + 1] + newf[i][j - 1] +
    19    4608000        1046318.9     0.2    38.4                  newf[i + 1][j] + newf[i - 1][j])
    20      1000           228.0     0.2     0.0      return newf

```

### memory\_profiler

```

Line #      Mem usage      Increment   Occurrences   Line Contents
=====
    13      81.254 MiB      31.961 MiB         1000      @profile
    14                               def gauss_seidel(f):
    15      81.254 MiB      -49.254 MiB         1000      newf = f.copy()
    16      81.254 MiB     -2415.875 MiB       49000      for i in range(1, len(newf) - 1):
    17      81.254 MiB     -115960.996 MiB    2352000      for j in range(1, len(newf[0]) - 1):
    18      81.254 MiB     -340783.230 MiB    6912000      newf[i][j] = 0.25 * (newf[i][j + 1] + newf[i][j - 1] +
    19      81.254 MiB     -227188.820 MiB    4608000      newf[i + 1][j] + newf[i - 1][j])
    20      81.254 MiB      -49.340 MiB         1000      return newf

```

It can be seen that the code in loop takes the most time.

**Task 2.3:** Use the Cython Annotation tool to identify the parts to use Cython

```

...
+06: def gauss_seidel(f):
+07:     newf = f.copy()
    08:
+09:     for i in range(1, newf.shape[0] - 1):
+10:         for j in range(1, newf.shape[1] - 1):
+11:             newf[i, j] = 0.25 * (newf[i, j + 1] + newf[i, j - 1] +
+12:                 newf[i + 1, j] + newf[i - 1, j])
    13:
+14:     return newf
...

```

It's too yellow here, so all the code need to be update.

**Task 2.4:** Use Cython to optimize the part you identified as the most computationally expensive. Compare the performance with the results obtained in Task 2.1.

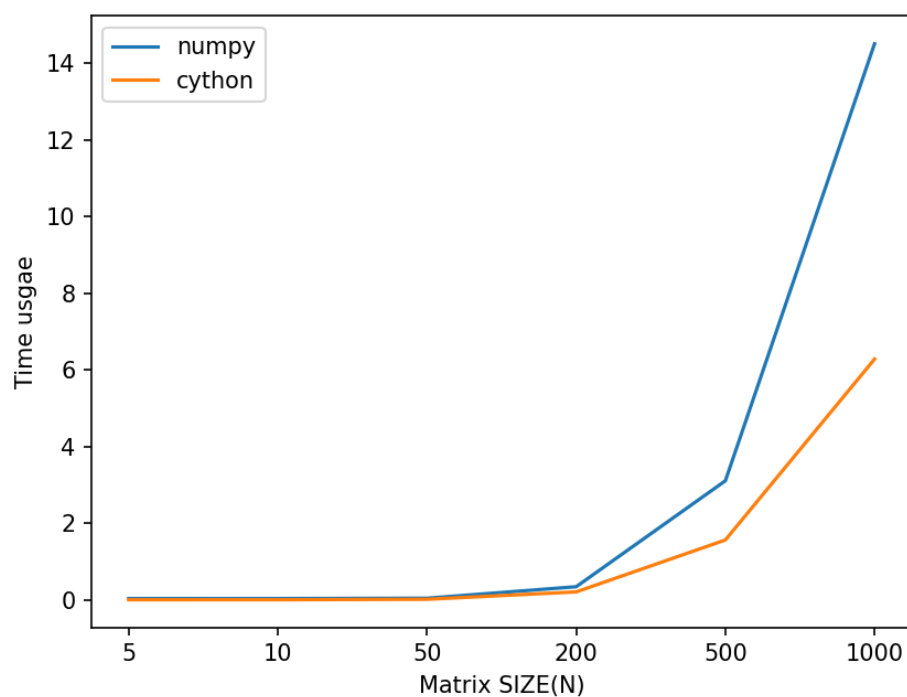
**After optimize:**

```

+01: #cython: language_level=3
02:
+03: def gauss_seidel(float[:, :] f, int N):
+04:     cdef float[:, :] newf = f.copy()
05:     cdef int i, j    # warning if not declare here: gauss_seidel.p
+06:     for i in range(1, N - 1):
+07:         for j in range(1, N - 1):
+08:             newf[i, j] = 0.25 * (newf[i, j + 1] + newf[i, j - 1] +
+09:                                   newf[i + 1, j] + newf[i - 1, j])
10:
+11:     return newf

```

Compare performance:



```

Numpy time   average:30091.61 ms    min:316.78 ms    max:145046.73 ms
Cython time  average:13433.94 ms    min:8.53 ms     max:62791.82 ms

-----FLOPS/s-----
+-----+-----+-----+
|      |      numpy      |      Cython      |
+-----+-----+-----+
|  5   | 7795.301615810118 | 292740.0468384001 |
| 10   | 63130.71530256975 | 1900057.0017099248 |
| 50   | 6112095.837662742 | 19530639.66751049 |
| 200  | 46955121.76196747 | 77659196.49912347 |
| 500  | 80517813.27790272 | 160209446.29500592 |
| 1000 | 137886552.28099278 | 318512581.915862  |
+-----+-----+-----+

```

It looks better after cython is applied, taking about half time compared with numpy.



**Task 2.5:** Use PyTorch to port your code to Nvidia GPUs. For this you will need to express the two nested loops operations as numpy roll operations in 2D as we did for the diffusion code.

**Task 2.6:** Use CuPy to port your code to Nvidia GPUs. See [C.3 Tutorial - Introduction to CuPy](#)

**Task 2.7:** Measure the performance (execution time) with GPU (PyTorch and cupy) and make a plot of the execution time varying the size the of the grid. Compare and comment of the performance difference with and without GPU.

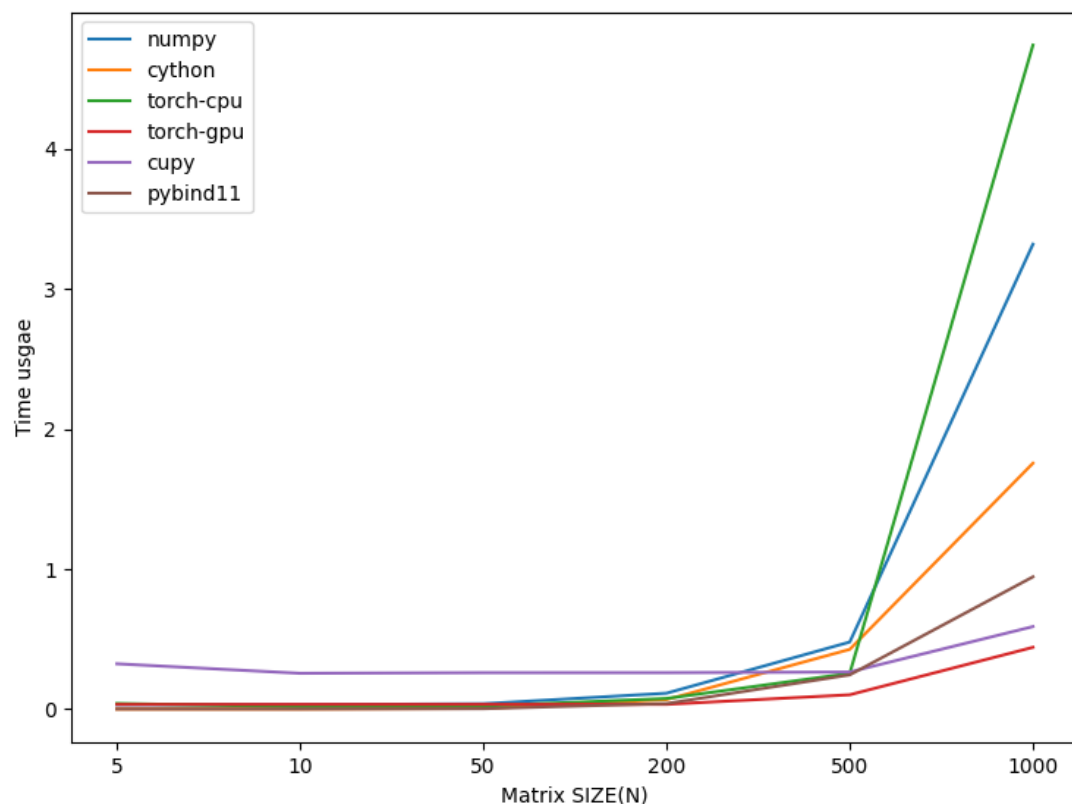
**Task 2.8:** Save the newgrid matrix as an hdf5 file using h5py (see the tutorial [C.4 Tutorial - The HDF5 Library/Module](#))

Its saved as file Assignment 3/newfs.hdf5 in github.

**Bonus Task B.1:** Develop the function you want to optimize in C/C++ or Fortran, make it a library, and couple C and Python. You can use the approach of your choice (see the lecture on foreign function interfaces). This is part of the bonus exercise.

**2.5-2.7 & B.1:** (codes in github Assignment 3/e2GaussSeidel\_gpu.py)

Set grid size as [5,10,50,200,500,1000]



Numpy time	average:6691.56 ms	min:311.78 ms	max:33210.36 ms				
Cython time	average:3762.01 ms	min:6.23 ms	max:17565.00 ms				
torch time	average:8598.33 ms	min:175.52 ms	max:47430.96 ms				
torch time	average:1131.07 ms	min:331.00 ms	max:4414.00 ms				
cupy time	average:3257.80 ms	min:2558.38 ms	max:5894.48 ms				
pybind11 time	average:2056.57 ms	min:8.66 ms	max:9448.77 ms				

-----FLOPS/s-----							
		numpy	Cython	torch-cpu	torch-gpu	cupy	pybind11
5	7834.7597767167745	488128.0489931737	6227.006777468889	7552.053738720625	772.128965995529	287999.5207529108	
10	64142.79931820587	2608759.1699805608	113928.85655562289	59219.74959700754	7817.354454971491	2086284.5566947178	
50	6539761.633020976	51854923.20019614	11389705.328249888	7393722.558022292	960794.3669502253	79884506.1783074	
200	140934861.80952805	235799866.7495115	218882646.41696766	467845586.8624004	61465712.84838072	403258022.37767655	
500	521844765.2471712	585960167.9606569	978101167.7579098	2442111934.4951224	942993319.1750535	1022711396.9922781	
1000	602228955.4731907	1138625289.8691829	421665034.350375	4530994555.978028	3392991781.334417	2116664103.4217467	

As the plots shows:

- 1、Performance with GPU is highly better when Matrix Size is big enough and time usage is stable as Matrix Size increasing. Because of the advantage of Synchronous Computing of GPU.
- 2、With GPU performance of torch is better than cupy.
- 3、Without GPU performance is: pybind11>cython>numpy. In particular, torch in cpu is better than cython when Matrix Size is small, and no better than numpy as Matrix Size big enough.

**B1:** Created two libs by Cython and Pybind11, Assignment3/gauss\_seidel\_cython.pyx, Assignment3/gauss\_seidel\_c.cpp, respectively.

## Bonus Exercise - Get familiar with VTK and Paraview

In this exercise, we ask you to experiment with one of the most used tools to visualize datasets coming from HPC simulations, Paraview. One of the most common and most used formats for visualization is called VTK.

As part of the exercise, you will need to repeat all the different steps presented in the [C.6 Tutorial - VTK & Visualization with Paraview](#)  
[C.6 Tutorial - VTK & Visualization with Paraview](#)

To install sputniPIC follow the instructions in the tutorial for paraview:

1. Install the **mpich** and **hdf5** libraries using **spack**
2. Install the SputniPIC code
3. When you run sputniPIC, the code will create a number of VTK files in the data folder (be sure to create a directory called data otherwise, you will get an error and you can't run sputniPIC)
4. Download and install paraview on your system
5. Use paraview for the visualization

**Task 3.1:** Make a volume plot and a slice of the rhoe at the last step recorded in the vtk file as shown in the tutorial. Add the snapshots you took to your report.

Get error when make sputniPIC:

```
(base) daybeha@daybeha:~/Documents/KTH/DD2358_HPC/HPC_assignment$ . /home/daybeha/Documents/KTH/DD2358_HPC/HPC_assignment/spack/share/spack/setup-env.sh
(base) daybeha@daybeha:~/Documents/KTH/DD2358_HPC/HPC_assignment$ spack load mpich
(base) daybeha@daybeha:~/Documents/KTH/DD2358_HPC/HPC_assignment$ spack load hdf5
(base) daybeha@daybeha:~/Documents/KTH/DD2358_HPC/HPC_assignment/sputniPIC$ make VERSION=CPU
mpicxx -std=c++11 -I./include -O3 -g -fopenmp -Wall src/RW_IO.cpp -c -o src/RW_IO.o
In file included from src/RW_IO.cpp:6:
./include/RW_IO.h:13:10: fatal error: hdf5.h: No such file or directory
   13 | #include <hdf5.h>
      |          ^~~~~~
compilation terminated.
make: *** [Makefile:60: src/RW_IO.o] Error 1
```

And when I'm trying to configure hdf5 from the source, I got weird output:

```
(base) daybeha@daybeha:~/Documents/KTH/DD2358_HPC/HPC_assignment/CMake-hdf5-1.14.0$ sh build-unix.sh
CMake Error at /home/daybeha/Documents/KTH/DD2358_HPC/HPC_assignment/CMake-hdf5-1.14.0/HDF5config.cmake:18 (cmake_minimum_required):
  CMake 3.18 or higher is required.  You are running version 3.16.3

(base) daybeha@daybeha:~/Documents/KTH/DD2358_HPC/HPC_assignment/CMake-hdf5-1.14.0$ cmake --version
cmake version 3.20.6

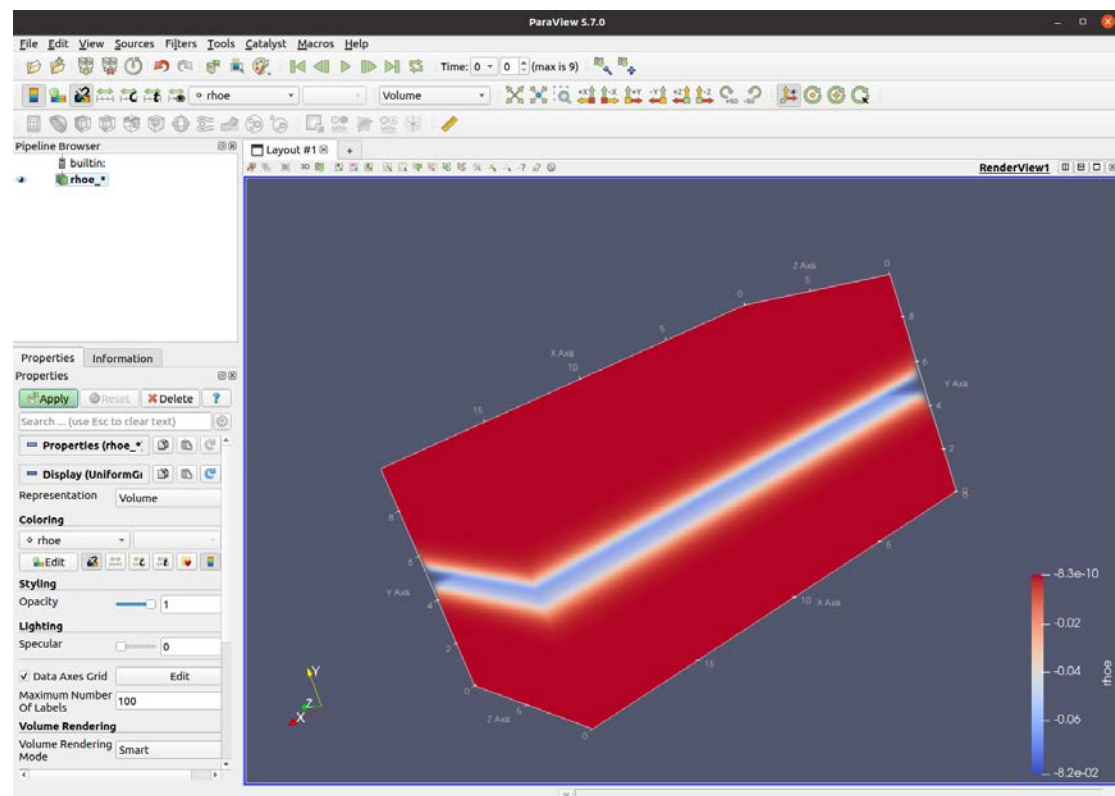
CMake suite maintained and supported by Kitware (kitware.com/cmake).
```

.....

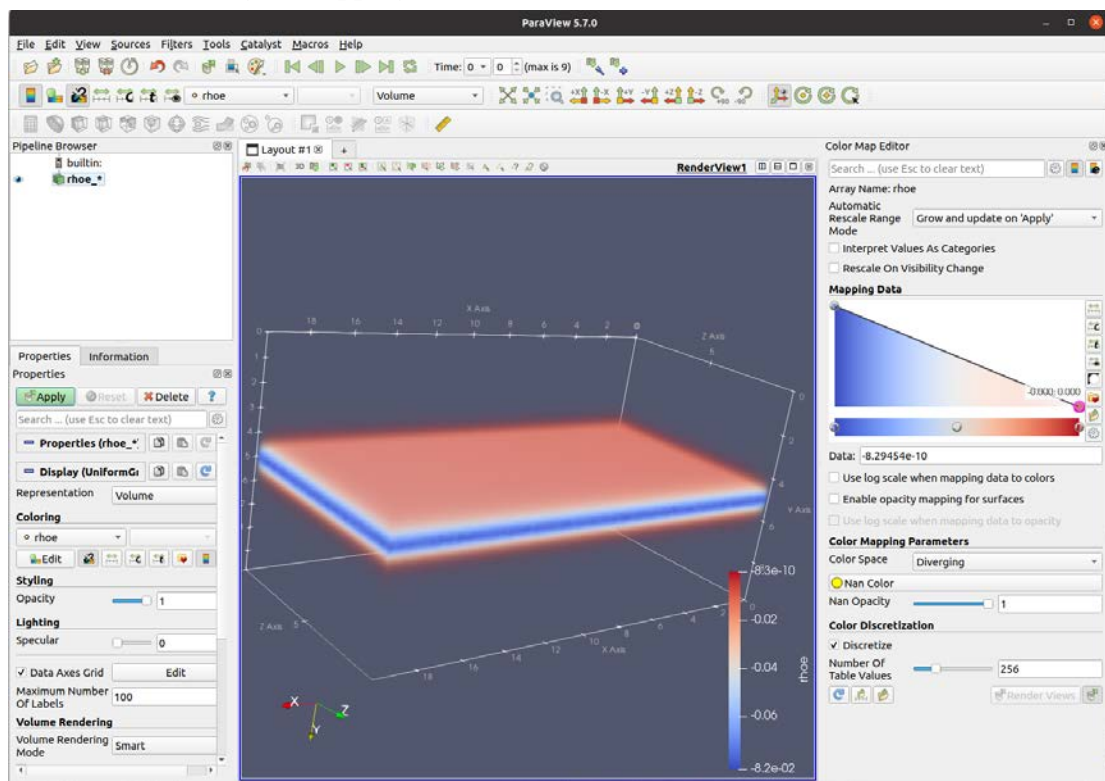
Lost of problems to load [hdf5](#) and [mpich](#) with spack using Ubuntu20.04.

Solved them by install hdf5 and mpich from source and build with cmake.

Successfully open rhoe\_\*.vtk with ParaView using **Volume rendering**:




reverse the coloring of Mapping Data:




play all the simulation output time step:

time=0

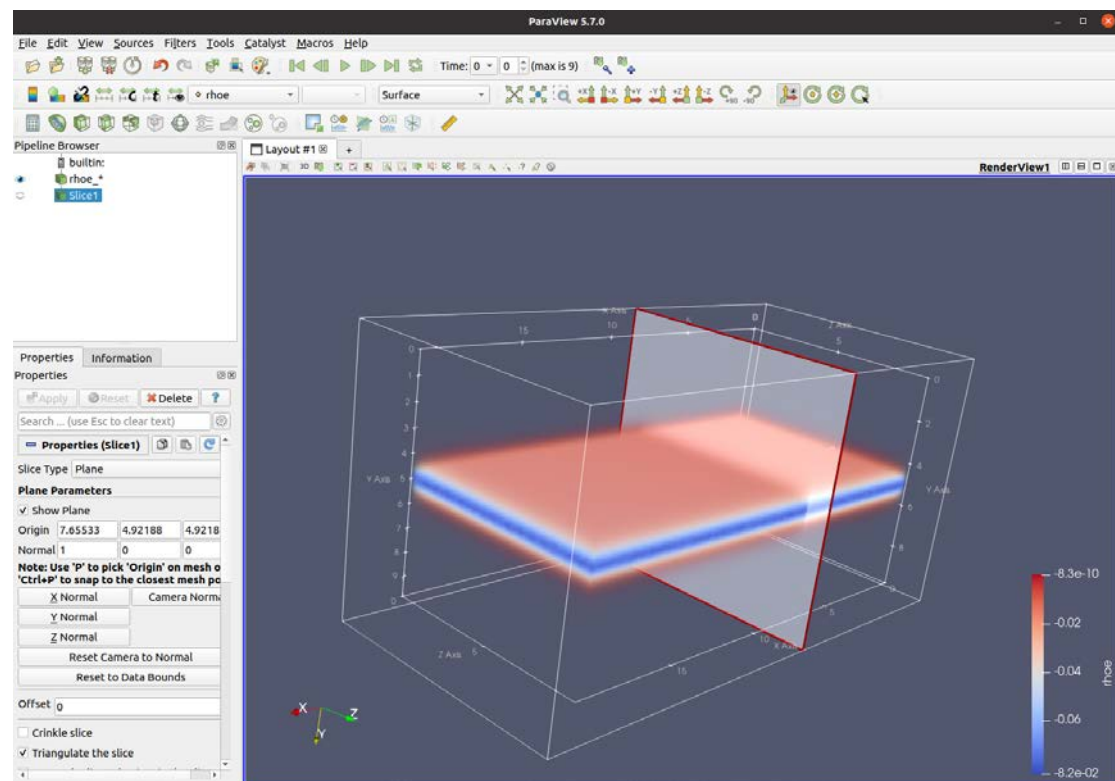


time=9 (max is 9)

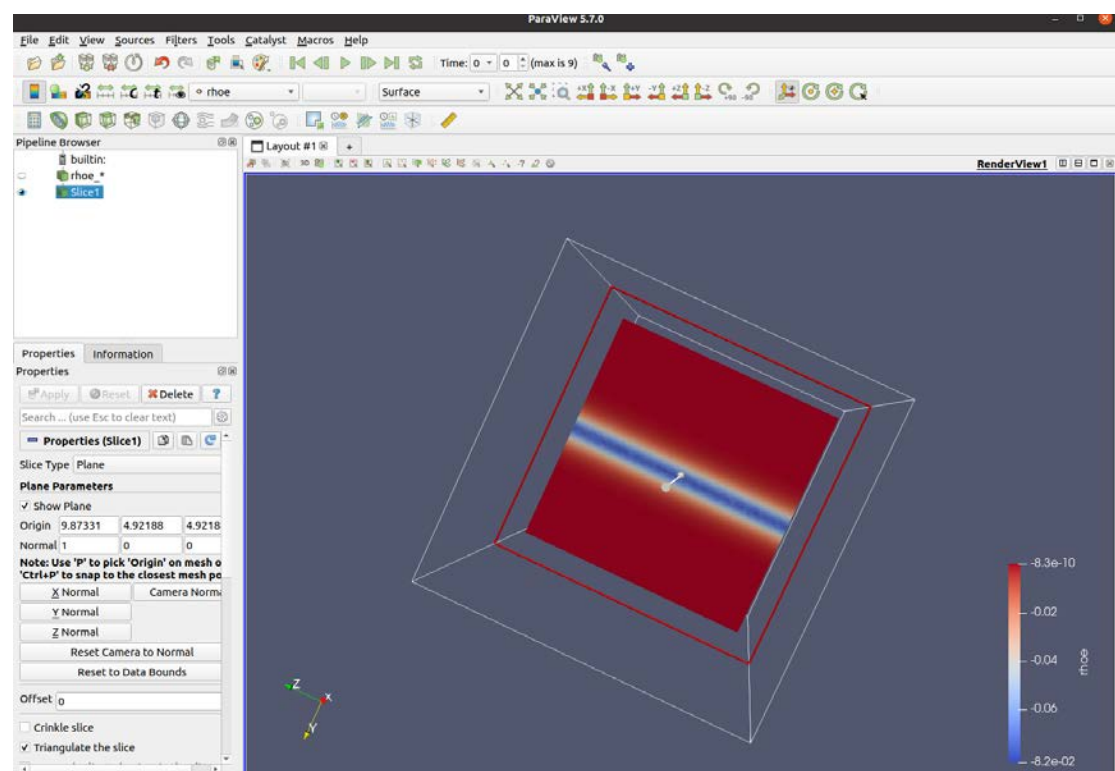


It's different in the middle of volume. Seems clearer.

To slice data:



data sliced:



show with rainbow desaturated:



